

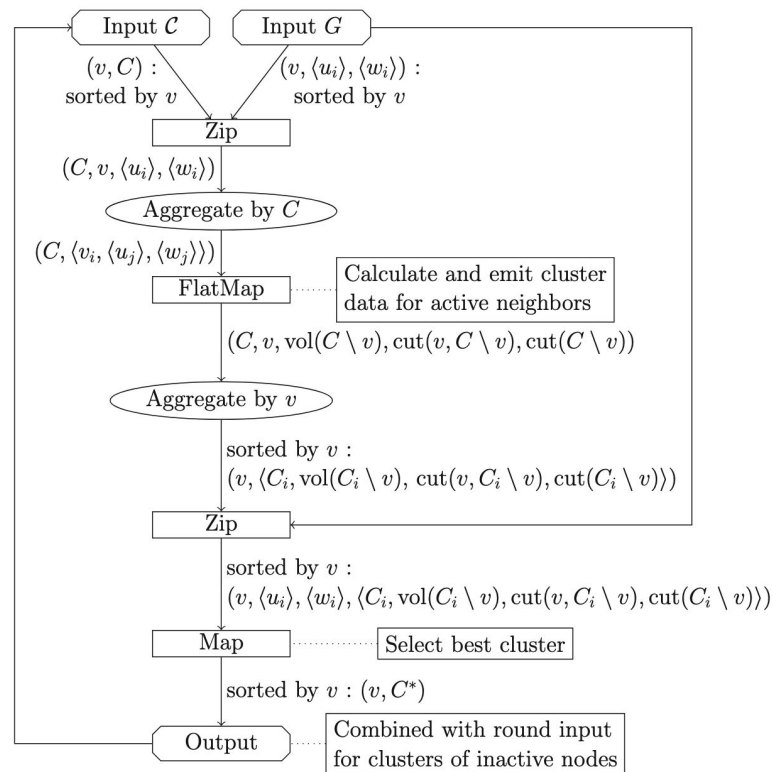
Distributing the Louvain Algorithm Using a Distributed Framework

Finn Archinuk
Ezra MacDonald
Alison Ziesel

April 06 2023
Code: github.com/aziesel/CSC502_Project

Learning Objectives


- Understand Louvain Algorithm
- Translate C++ implementation to PySpark
- Optimize implementation



Hamann *et al.* 2018

Communities in Networks

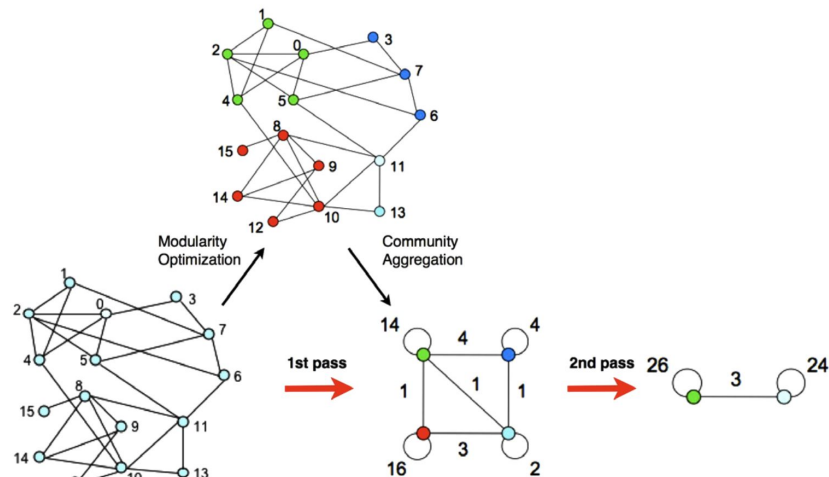
- What are communities?
 - Communities are a subset of nodes that are densely connected to each other and loosely connected to other nodes in the graph
- Applications
 - Network analysis
 - Recommender systems
 - similarity based algorithms
 - Any problem where you have a graph
- How are they defined?
 - Weights and graph edges
 - Edges with a greater weight in general indicate a closer relationship between two vertices, holding everything else constant



Fortunato & Castellano 2007

Louvain Algorithm

- What does it do?
 - divides network into most modular, or closely related communities by the given weighted edge metric
- Why is it useful?
 - It's fast
 - Has properties that can be exploited to increase parallelizability
 - It's relatively simple to implement
 - It's unsupervised and outperforms many similar modularity optimization methods [Aynaud et. al 2011]
- How does it work?
 - Iteratively maximize the modularity score until:
 - i. A fixed number of rounds has been reached
 - ii. No moves to another cluster can be made in a given round



Blondel *et al.* 2008

Network Modularity

- Σ_{in} = Sum of all the edge weights between nodes within the community C
- Σ_{tot} = Sum of all the edge weights for nodes within the community including edges connected to other communities

$$Q_c = \frac{\Sigma_{in}}{2m} - \left(\frac{\Sigma_{tot}}{2m}\right)^2$$

Blondel et al. 2008

- In general higher modularity scores = better cluster
- $\text{vol}(C)$ = sum of weights connecting each node within a cluster
- $\text{cut}(C)$ = Sum of weights of all vertices $v \in C$ for pairs (u, v) such that $v \in C$ and $u \notin C$

$$Q(\mathcal{C}) := \sum_{C \in \mathcal{C}} \frac{\text{vol}(C) - \text{cut}(C)}{\text{vol}(V)} - \sum_{C \in \mathcal{C}} \frac{\text{vol}(C)^2}{\text{vol}(V)^2}.$$

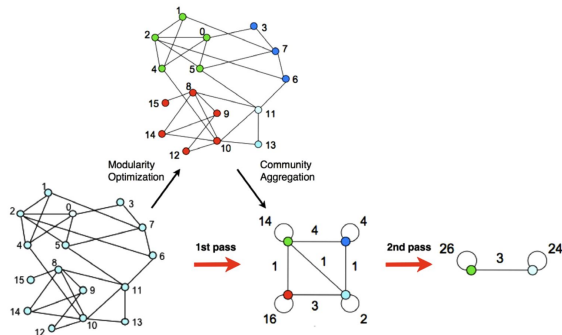
Hamann et al. 2018

Louvain Algorithm

Algorithm LouvainAlgorithm(Graph G)

Require: $(x_1, y_1), \dots, (x_m, y_m)$ where $x_i \in X$ and $y_i \in \{-1, 1\}$

- 1: $C \leftarrow \text{index}(G)$ Map the index of each node to its own cluster
- 2: $G' = G$
- 3: $q \leftarrow -\infty$
- 4: **while** $q < Q(G', C)$ **do**
- 5: $q \leftarrow Q(G', C)$
- 6: $C \leftarrow \text{MoveNodes}(G')$ // Phase 1
- 7: $G' \leftarrow \text{Aggregate}(G', C)$ // Phase 2
- 8: $C \leftarrow$ each node of G' in its own community
- 9: **end while**
- return** G'



Algorithm MoveNodes(Graph G)

Require: C the index of communities for each nodes of G

Require: $(x_1, y_1), \dots, (x_m, y_m)$ where $x_i \in X$ and $y_i \in \{-1, 1\}$

- 1: $C \leftarrow \text{index}(G)$ Map the index of each node to its own cluster
- 2: $G' = G$
- 3: $q \leftarrow -\infty$
- 4: **while** at least one node moves or reached max iterations **do**
- 5: **for** random $v \in V(G)$ **do**
- 6: $\text{max}Q \leftarrow -\infty$
- 7: $\text{max}C \leftarrow$ community of v
- 8: **for** each neighbor u of v **do**
- 9: $\text{delta}Q = \Delta Q$ between v and n
- 10: **if** $\text{max}Q < \text{delta}Q$ **then**
- 11: $\text{max}Q \leftarrow \text{delta}Q$
- 12: $\text{max}C \leftarrow$ community of u
- 13: **end if**
- 14: **end for**
- 15: $C \leftarrow \text{Update}(C, v)$ // Update cluster for node v
- 16: **end for**
- 17: **end while**
- return** G'

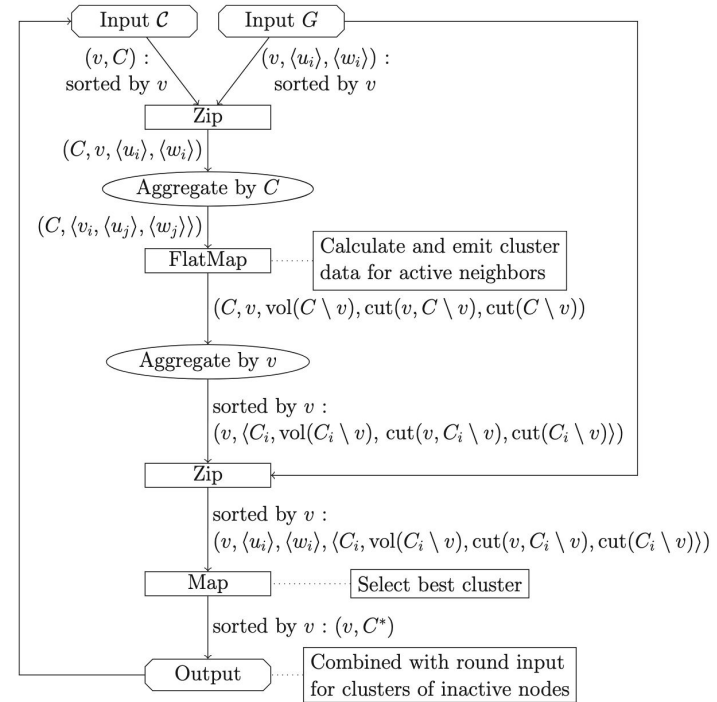
Algorithm Aggregate(Graph G, Partition C)

- 1: $G' \leftarrow$ aggregate nodes in the same community based on C

return G'

Distributed Clustering Using Modularity and Map Eq.

- Group graph inputs with the individual vertices and their corresponding cluster assignments
- Generate cluster groupings based on vertices that are randomly assigned to move in this round
- Compute the volumes and cuts for each generated clustering
- Sort respective clusterings and update



Opportunities for Parallelization

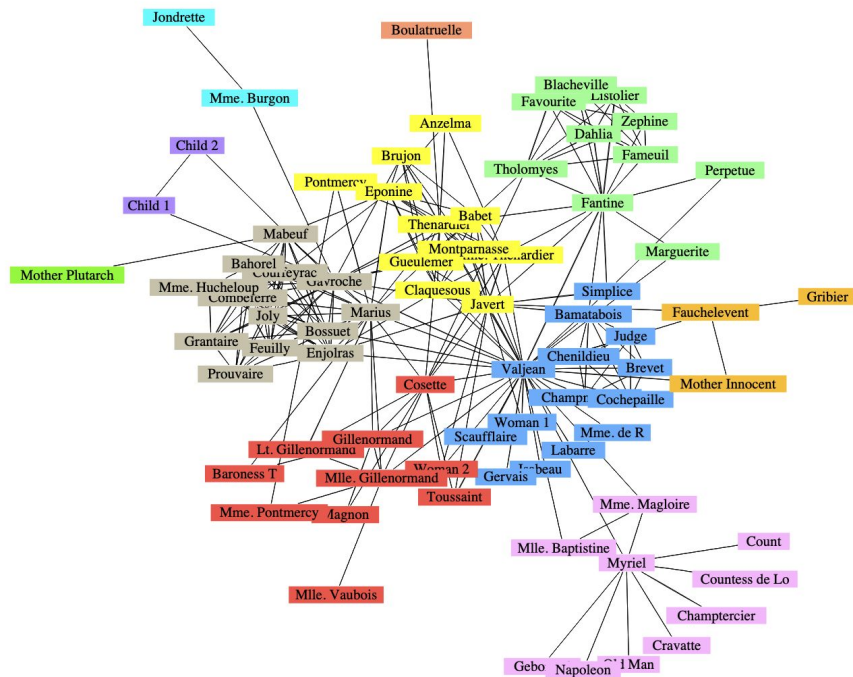
- Evaluation of the modularity gain for all vertices (Computing ΔQ)
 - Each vertex computes the modularity gain for joining its neighbours' communities
 -
 - This results in a considerable amount of
- Clustering that's required at each round requires multiple sort operations
-

Why PySpark

- The paper by Hamann et. al which we chose to focus on for this project drew many parallels with the MapReduce model
- Non-parallel implementations in Python exist
- Python integration into data science
 - Data science has lots of interest in network/community analysis
- No currently publicly available PySpark implementation of the Louvain algorithm
- Other implementations: Spark, Thrill (Hamann 2018)
 - Other implementations are thesis work...

Experimentation: *Les Misérables* dataset

- Toy example: Les Misérables character interaction network
 - Edges indicated the count of co-appearances of characters
- Used for development
- No real ground truth







Knuth 1993

Results: Les Miserables

- Took a collaborative approach to development
- Prioritized having at least one working implementation
- Some implementations work better than others
- Gave everyone more development time
- More opportunity to explore unique implementation approaches

Two implementations:

- 1) 52/77 (70.1%) concurrence
- 2) 60/77 (77.9%) concurrence



Experimentation: ArXiv dataset

- A collection of 2.2 million+ metadata records of articles published on ArXiv retrieved from [Kaggle](#).
- Explored the connection between coauthors
 - Edges were created based on co-authoring a publication
 - Edge weights were computed based on the number of shared publications
- No ground truth has been established for the dataset... unsupervised

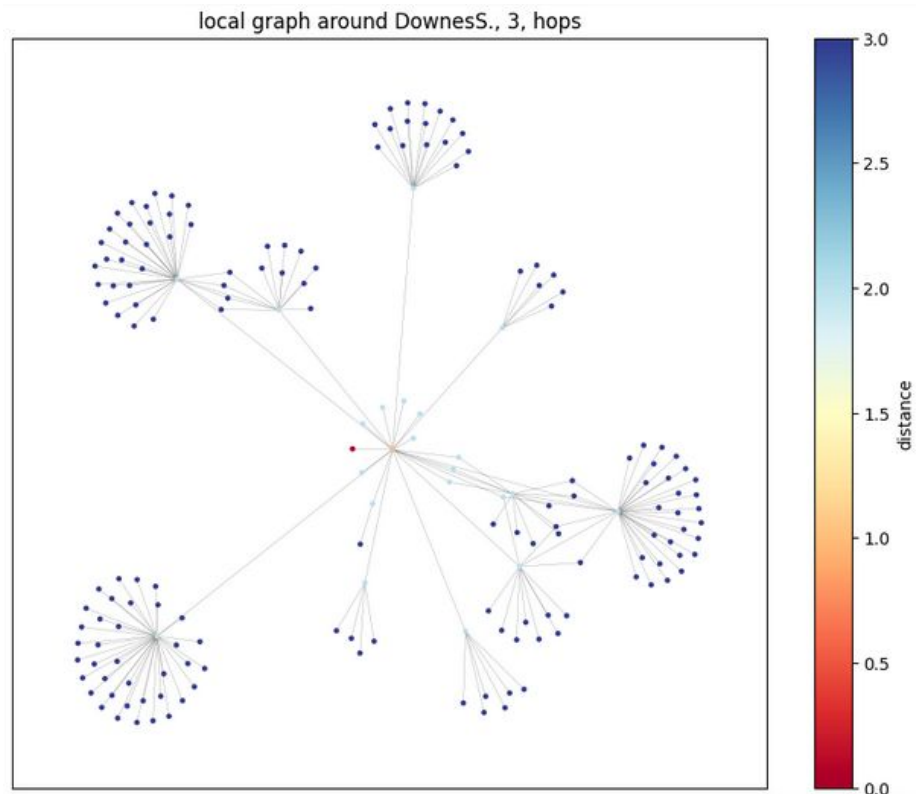
```
▼ "root" : { 14 items
  "id" : string "0704.0001"
  "submitter" : string "Pavel Nadolsky"
  "authors" : string "C. Bal\`azs, E. L. Berger, P. M. Nadolsky, C.-P. Yuan"
  "title" :
    string "Calculation of prompt diphoton production cross sections at
    Tevatron and LHC energies"
  "comments" : string "37 pages, 15 figures; published version"
  "journal-ref" : string "Phys.Rev.D76:013009,2007"
  "doi" : string "10.1103/PhysRevD.76.013009"
  "report-no" : string "ANL-HEP-PR-07-12"
  "categories" : string "hep-ph"
  "license" : NULL
  "abstract" :
    string " A fully differential calculation in perturbative quantum
    chromodynamics is presented for the production of massive photon pairs at
    hadron colliders. All next-to-leading order perturbative contributions
    from quark-antiquark, gluon-(anti)quark, and gluon-gluon subprocesses are
    included, as well as all-orders resummation of initial-state gluon
    radiation valid at next-to-next-to-leading logarithmic accuracy. The
    region of phase space is specified in which the calculation is most
    reliable. Good agreement is demonstrated with data from the Fermilab
    Tevatron, and predictions are made for more detailed tests with CDF and
    DO data. Predictions are shown for distributions of diphoton pairs
    produced at the energy of the Large Hadron Collider (LHC). Distributions
    of the diphoton pairs from the decay of a Higgs boson are contrasted with
    those produced from QCD processes at the LHC, showing that enhanced
    sensitivity to the signal can be obtained with judicious selection of
    events. "
  ▼ "versions" : [ 2 items
    ▶ 0 : {...} 2 items
    ▶ 1 : {...} 2 items
```

Results: Arxiv Figures


Visualization: “DownesS.” only has 1 direct connection. This figure is of the local neighbourhood.

Preprocessing: Removed “Collaborations”

The next slide is a continuation to 4 hops.
“Fireworks” are highly connected individuals.




local graph around DownesS., 4, hops




Authors were compacted.
These are a random set of
5000 *edges* of those compacted clusters

Highly connected authors make
this very difficult.



First 5000 compacted edges


Some obvious structure



Authors were clustered starting from 100k singletons to ~15k clusters. These clusters were then compacted.

The compacted clusters went through a single iteration (10 subrounds). This is the resulting cluster and neighbours of a *single* cluster.

Because of highly connected authors, compaction isn't very useful. Another challenge is highly connected authors make small communities hard to distinguish.



1e6

5

4

3

2

1

17

Conclusions

- The Louvain algorithm is an easily parallelizable algorithm that returns quality results and can be fast
- PySpark can be used to improve the computational efficiency of the Louvain algorithm
- We got acceptably close results of the non-distributed implementation
- The ArXiv dataset was a worthy adversary

References

Aynaud, T.; Blondel, V. D.; Guillaume, J.; Lambiotte, R. "Multilevel local optimization of modularity". *Graph Partitioning*. John Wiley & Sons: 315–345. (2011)

Blondel, V.D. *et al.* "Fast unfolding of communities in large networks." *Journal of Statistical Mechanics: Theory and Experiment* (2008)

Hamann, M. *et al.* "Distributed graph clustering using modularity and map equation." *Euro-Par 2018: Parallel Processing: 24th International Conference on Parallel and Distributed Computing* (2018)

<https://github.com/taynaud/python-louvain>

Hagberg, A.A. *et al.* "Exploring network structure, dynamics and function using NetworkX" *Proceedings of the 7th Python in Science Conference (SciPy2008)* (2008)

